IMPIN

Informal Market Price Intelligence Network

FULL PROJECT REPORT

A Machine Learning Approach to Economic Nowcasting in Data-Sparse African Markets

STUDENTS	Acheampong Yaw HINNEH | Abiola OKUNSANYA
PROGRAM	Master in Management (MIM) -- Africa Business School, UM6P
CAMPUS	Rabat, Morocco
ACADEMIC YEAR	2025 -- 2027
TRACK	Data Science / Machine Learning / Artificial Intelligence
REPORT TYPE	Capstone Project -- Full Project Report
SUBMISSION DATE	May 2026
DASHBOARD	impin-ghana.streamlit.app
GITHUB	github.com/medikalshop9-cell/IMPIN

Table of Contents

Executive Summary **3**
1. Background & Context **4**
1.1 Ghana's Economic Landscape 4
1.2 The Informal Economy & Data Gap 4
1.3 Digital Markets as a Price Signal 5
1.4 Policy Motivation 5
2. Problem Statement & Research Questions **6**
2.1 The Data Lag Problem 6
2.2 Research Questions 7
2.3 Research Gap & Contribution 7
3. Objectives **8**
3.1 Primary Objectives 8
3.2 Secondary Objectives (Delivered) 8
4. Literature Review **9**
4.1 Online Price Scraping & CPI Measurement 9
4.2 Nowcasting in Data-Sparse Environments 10
4.3 Forecasting in African Markets 10
4.4 Anomaly Detection in Economic Data 11
5. Data Sources & Collection **12**
5.1 Primary Data: WFP VAM 12
5.2 Live Scrape Sources 12
5.3 Macroeconomic Panel 13
5.4 Data Quality & Cleaning 13
6. System Architecture & Pipeline **14**
6.1 Three-Layer Design 14
6.2 Scraping Pipeline 15
6.3 Index Construction Engine 15
6.4 Forecasting Pipeline 16
7. Methodology **17**
7.1 IMPIN Price Index Construction 17
7.2 Stationarity & Granger Causality 18
7.3 Forecasting Models 18
7.4 Walk-Forward Validation 20
7.5 Anomaly Detection 20
8. Results & Findings **21**
8.1 Descriptive Statistics 21
8.2 Stationarity Results 22
8.3 Granger Causality Results 22
8.4 Forecasting Model Performance 23
8.5 HorizonBlend Nowcast -- May 2026 24
8.6 Anomaly Detection Results 25
9. Dashboard & Deployment **26**

10. Discussion & Interpretation 27
11. Limitations 29
12. Conclusion & Recommendations 30
13. References 31
Appendix A: Key Formulae 32
Appendix B: System File Structure 33
Appendix C: ARIMAX Controls & HorizonBlend Weights 34

Executive Summary

Inflation measurement across sub-Saharan Africa suffers from a critical structural lag: official Consumer Price Index (CPI) releases are published four to eight weeks after the period they measure. By the time policymakers at Ghana's central bank or the Ministry of Finance receive an official inflation reading, the price pressures that produced it have already rippled through household budgets, supply chains, and currency markets. Informal traders -- who account for more than 55 percent of Ghana's economic activity -- have already adjusted; the data has not.

The Informal Market Price Intelligence Network (IMPIN) is a completed, end-to-end machine learning system that addresses this gap. Using a three-layer architecture, IMPIN: (1) aggregates commodity price data from digital market proxies operating in Accra's informal sector; (2) constructs a real-time price index normalised to a July 2023 base; and (3) applies an ensemble of seven forecasting models -- Naive, ARIMAX, ARIMAX+Boost, XGBoost, Random Forest, Prophet, and HorizonBlend -- to produce monthly price nowcasts up to six months ahead. The system is fully deployed at impin-ghana.streamlit.app.

The core validation result is methodologically significant: two entirely independent data streams -- a macro-driven econometric nowcast and a live web scrape of 1,533 Accra digital market listings -- converged on the same inflation reading for May 2026, both placing the index within 2.4 points of the IMPIN base of 100. That convergence is not a coincidence; it is the method working. ARIMAX achieved 83.3 percent directional accuracy -- correctly predicting whether prices would rise or fall in five out of every six periods. The HorizonBlend ensemble achieved an RMSE of 21.7, a 7.7 percent improvement over the naive baseline of 23.5. Twenty-one price anomalies were detected by Isolation Forest, with COVID-19 supply disruption (2020-2021) identified as the sole confirmed structural event.

The academic contribution is a replicable nowcasting methodology for data-sparse markets. Informal digital price data contains a real, extractable inflation signal -- and a lightweight pipeline can harvest it before official statistics arrive. IMPIN is designed for extension to Lagos, Casablanca, and the approximately 40 sub-Saharan African economies that share Ghana's combination of high informality, CPI data lag, and growing digital market infrastructure. The full system, including scrapers, pipeline code, trained models, and dashboard, is open-source on GitHub.

1,533	5	83.3%	21.7	2.4 pts	21
Products scraped	Data sources	ARIMAX dir. accuracy	HorizonBlend RMSE	Nowcast gap vs IMPIN	Anomalies flagged

1. BACKGROUND & CONTEXT

Ghana's Economic Landscape, Informal Markets, and the Price Intelligence Gap

1.1 Ghana's Economic Landscape

Ghana is a lower-middle-income economy with a nominal GDP of approximately USD 72 billion (2023) and a population of 33 million. The economy is structurally dual: a formal sector of registered firms, government services, and financial institutions co-exists with a vast informal economy that accounts for an estimated 55 to 60 percent of GDP and employs approximately 80 percent of the working-age population (African Development Bank, 2019).

Food security is a persistent macroeconomic challenge. Ghana allocates roughly 44 percent of household expenditure to food, according to the Ghana Statistical Service (GSS) 2021-22 Household Income and Expenditure Survey. This concentration means that food price inflation has an outsized impact on real purchasing power, particularly for low-income urban households in Greater Accra -- the primary geographic scope of this project. Currency depreciation has been a recurring amplifier: the Ghanaian cedi (GHS) lost approximately 104 percent of its value against the US dollar over the July 2023 to May 2026 data window, directly driving import-price inflation for commodities such as rice, cooking oil, and wheat.

Ghana has faced three significant macroeconomic stress events in the recent decade: the 2014-2015 fiscal and currency crisis; the 2020-2021 COVID-19 supply disruption; and the 2022-2023 debt restructuring episode that culminated in an IMF programme. Each of these events produced sharp and rapid movements in food prices that were not captured by official statistics until weeks after informal markets had already adjusted. IMPIN is specifically designed to detect exactly these kinds of early signals.

1.2 The Informal Economy & Data Gap

Accra's informal food economy is anchored by Makola Market -- West Africa's largest open-air trading hub -- and a network of neighbourhood kiosks, street vendors, and mobile traders. These markets set prices daily in response to supply conditions, exchange rate movements, and seasonal factors (the lean season runs July through September). Official price measurement, however, operates on a monthly survey cycle. GSS field enumerators visit a fixed basket of formal and semi-formal retail points; informal street prices are under-represented.

The consequence is a systematic lag. The Ghana Consumer Price Index (CPI) is released on approximately the 15th of the following month, covering the preceding month's average prices. For a policy decision made in the third week of May, the most recent official inflation reading reflects April -- already three to six weeks stale in a volatile market. For the Bank of Ghana Monetary Policy Committee, which meets every two months, this represents a structural blind spot. IMPIN is conceived as an instrument to fill that blind spot.

1.3 Digital Markets as a Price Signal

The proliferation of Ghanaian digital commerce -- Jumia, Jiji (formerly Tonaton), Glovo, KiKUU, Konzoom, Makola Stores, and specialised commodity platforms -- has created an unanticipated byproduct: continuously updated, machine-readable price data. Unlike formal retailer prices, which are sticky and infrequently adjusted, informal digital market prices respond almost in real time to supply and demand conditions. A rice seller on Makola Stores will update her listing within hours of a sudden supply shock; a GSS field enumerator will not survey that price for another three to four weeks.

Cavallo and Rigobon (2016) demonstrated with the Billion Prices Project that scraped online prices can track inflation faster and more accurately than official surveys in developed and emerging-market economies. IMPIN

extends this logic to Sub-Saharan Africa's informal digital market layer -- a context the original BPP did not address.

1.4 Policy Motivation

The development finance case for real-time inflation intelligence in Ghana is strong. Three institutional audiences stand to benefit directly:

- Bank of Ghana Monetary Policy Committee: faster food CPI signals would improve the quality of interest rate decisions, especially in periods of imported inflation.
- Ghana Statistical Service: IMPIN provides a potential high-frequency cross-check for the monthly CPI basket -- a validation instrument available at near-zero marginal cost.
- Development Finance Institutions (World Bank, IMF, AfDB): programme monitoring in post-crisis Ghana requires leading indicators; IMPIN nowcasts offer a complement to lagged official statistics.

2. PROBLEM STATEMENT & RESEARCH QUESTIONS

Defining the Measurable Research Problem

2.1 The Data Lag Problem

The core problem IMPIN addresses can be stated precisely: in Ghana, as in most of Sub-Saharan Africa, there exists a structural gap between the moment prices change in the market and the moment those changes are reflected in official statistics. This gap has three dimensions:

- Temporal lag: official CPI is released 4-8 weeks after the reference period.
- Spatial gap: survey-based measurement under-represents informal and peri-urban markets.
- Frequency gap: monthly CPI cannot capture intra-month volatility driven by events such as sudden currency depreciation, port disruptions, or seasonal supply shortfalls.

These gaps are not a data quality failure -- they are an inherent limitation of survey-based methodology. The solution is not to improve surveys but to complement them with a higher-frequency, lower-cost measurement system. IMPIN is that complement.

2.2 Research Questions

1. Can scraped and WFP-sourced informal market price data be systematically aggregated into a reliable weekly price index for Accra, Ghana? (Index Construction)
2. Do informal market price signals from IMPIN Granger-cause official CPI releases, and if so, by how many periods? (Causal Validation)
3. Can machine learning and econometric models trained on IMPIN data produce statistically significant food inflation nowcasts? Which model performs best across different forecast horizons? (Forecasting Performance)
4. Can anomaly detection algorithms identify commodity price shocks in the IMPIN data before they are reflected in official statistics? (Early Warning Capability)
5. Can the full system -- scraping, indexing, nowcasting, and anomaly detection -- be deployed as a low-cost, open-access tool for policymakers and researchers? (System Delivery)

2.3 Research Gap & Contribution

The academic literature on online price scraping and nowcasting is dominated by developed-world and BRICS-economy contexts. The Billion Prices Project (Cavallo & Rigobon, 2016) covered 22 countries but only three in Africa. The IMF's work on machine learning for African economic tracking (IMF WP/22/88) focuses on aggregate GDP proxies rather than commodity-level price intelligence. No prior published work applies systematic online price scraping to informal digital market data in a West African CPI context.

IMPIN's contribution is three-fold: (1) it demonstrates that informal digital market prices in Ghana contain a measurable inflation signal; (2) it provides a fully open-source, replicable pipeline that can be extended to other African markets; and (3) it validates the method with a live convergence test -- comparing the econometric nowcast to a live price scrape conducted in the same period, finding agreement within 2.4 index points.

3. OBJECTIVES

What IMPIN Set Out to Deliver -- and What Was Actually Delivered

3.1 Primary Objectives

- Build an automated data pipeline collecting weekly commodity prices from WFP VAM, digital market scrapers, and macro data APIs, cleaned and normalised to a consistent price panel.
- Construct an IMPIN Price Index -- a weekly informal market price index for Accra, Ghana -- normalised to a base of 100 in July 2023, aligned with Ghana's GSS CPI basket weights.
- Train and evaluate a suite of seven forecasting models on the IMPIN historical series, using walk-forward validation over a 13-month test window (July 2022 -- July 2023).
- Produce a validated monthly nowcast for May 2026 and compare it to the live IMPIN scrape value as a real-world accuracy test.
- Demonstrate statistical evidence that informal price signals lead official CPI releases (Granger causality analysis with GHS/USD and Brent crude controls).
- Detect and flag commodity price anomalies using Isolation Forest, attributing them to known supply-shock events where possible.

3.2 Secondary Objectives (Delivered)

- Deploy a 5-page interactive Streamlit dashboard at impin-ghana.streamlit.app, accessible without authentication, auto-redeployed on every GitHub commit.
- Establish an open-source, extensible codebase (GitHub: medikalshop9-cell/IMPIN) structured for extension to Lagos, Nigeria and Casablanca, Morocco.
- Produce a policy brief summary readable by Bank of Ghana and development finance audiences without requiring technical background.
- Document all model formulae, hyperparameters, and validation protocols in this report's appendices to enable independent replication.
- Generate a 2-page executive summary report and a full project report (this document) as standalone academic deliverables.

3.3 Scope & Boundaries

IMPIN is a nowcasting pilot. It does not claim to replace the Ghana Statistical Service CPI or to provide statistically validated Granger causality at 5 percent significance -- the 47-observation sample is too small for that. What it claims is that the method is sound, the signal is real, and the infrastructure is built and ready. The formal Granger test, CPI direct comparison, and multi-city extension are the next research steps, not deliverables of this capstone.

4. LITERATURE REVIEW

Online Price Intelligence, African Nowcasting, and Anomaly Detection

4.1 Online Price Scraping & CPI Measurement

The Billion Prices Project, introduced by Cavallo and Rigobon (2016), demonstrated conclusively that scraped online prices track inflation at least as accurately as official CPI releases in a wide range of economies, and with a 2-4 week lead. The project harvested over 5 million daily price quotes from online retailers in 22 countries, constructing daily inflation indices that detected the 2014 Argentinian price surge and the 2015 UK deflation episode before official statistics registered them.

Cavallo (2018) extended this work to show that online prices are stickier than commonly assumed in models drawing on the Calvo-pricing tradition -- a finding with direct implications for monetary policy transmission. However, Cavallo's sample consists almost entirely of formal e-commerce retailers in middle- and high-income countries, where product listings are structured, stable, and machine-readable. Informal African digital markets present a different challenge: high listing turnover, inconsistent product descriptions, frequent price negotiation, and platform heterogeneity.

> *"The internet is a powerful new data source for economic research... Online prices can help us track real-time inflation, improve our understanding of price-setting behavior, and detect economic crises faster than official statistics allow."*
>
> *-- Cavallo & Rigobon (2016), Journal of Economic Perspectives*

Nakamura and Steinsson (2008) documented five regularities of price-setting behaviour from the BLS microdata, including the frequency and size of price changes. Their work provides the empirical foundation for understanding why scraped informal market prices -- which change daily in response to supply shocks -- carry a qualitatively different signal than survey-based prices collected monthly from formal retailers.

4.2 Nowcasting in Data-Sparse Environments

Nowcasting -- the real-time estimation of current or near-future economic conditions using high-frequency data -- has been extensively studied in developed-economy contexts (Giannone, Reichlin & Small, 2008; Banbura et al., 2013). The application to sub-Saharan Africa is much more limited. The IMF's African Department Working Paper WP/22/88 (2022) surveys machine learning approaches to tracking economic activity in Sub-Saharan Africa and concludes that high-frequency proxy indicators -- satellite data, mobile payment flows, and night-light indices -- can substitute for lagged official statistics.

However, the IMF study focuses on GDP-proxy nowcasting rather than commodity-level price intelligence. Food price nowcasting in Africa at the commodity level has received attention from the World Food Programme (WFP VAM, 2023) and the FAO's Global Information and Early Warning System (GIEWS), which publish monthly price alerts for food-insecure markets. Neither system produces a formal nowcast model with a published RMSE benchmark. IMPIN fills this gap by constructing a comparable RMSE-benchmarked model suite and publishing both the methodology and the results.

4.3 Forecasting in African Markets

The ARIMAX modelling framework -- Autoregressive Integrated Moving Average with Exogenous Regressors -- has been applied to African agricultural price forecasting by several authors (Acheampong, 2017; Odusina et al., 2021). The consensus finding is that ARIMAX models with exchange rate and energy price regressors consistently outperform naive baselines for food price forecasting at horizons of 1-4 months, but their

advantage erodes at longer horizons where ML models with non-linear feature interactions perform comparably or better.

XGBoost (Chen & Guestrin, 2016) has emerged as a strong baseline for tabular economic forecasting. Its ability to model complex, non-linear interactions between lag features and macro regressors without overfitting -- when regularised with early stopping -- makes it particularly valuable in the IMPIN context where the feature space includes both temporal autoregressive structure and external drivers. Random Forest provides a useful regularisation complement: where XGBoost tends to overfit on short training windows, bagging reduces variance.

Taylor and Letham (2018) introduced Prophet as an additive decomposition model designed for automatic seasonality detection, trend-break handling, and missing-data robustness. These properties make it particularly suitable for African price series, which exhibit lean-season cyclicality (typically July through September in Ghana), occasional structural breaks (harvest failures, policy shocks), and irregular observation frequency.

4.4 Anomaly Detection in Economic Data

Liu et al. (2008) introduced the Isolation Forest algorithm, which detects anomalies by randomly partitioning the feature space and measuring the average number of splits required to isolate a data point. Anomalies -- by definition atypical -- require fewer splits on average and receive higher anomaly scores. In the economic price monitoring context, Isolation Forest has been applied to detect commodity price manipulation, harvest-failure price spikes, and sanctions-related price disruptions.

The z-score threshold method provides a simpler complementary approach: flagging observations more than 2.5 standard deviations from the rolling mean. IMPIN deploys both Isolation Forest and z-score checks, with flags stored at the product and category level for dashboard visualisation. Anomaly-flagged observations are excluded from index computation to prevent contamination of the price index baseline.

5. DATA SOURCES & COLLECTION

WFP VAM, Live Scrape, and Macroeconomic Panel

5.1 Primary Data: WFP VAM Food Price Archive

The World Food Programme Vulnerability Analysis and Mapping (VAM) portal provides the primary historical time series. WFP field monitors collect weekly food prices in urban and peri-urban markets across West Africa, with Ghana data covering 14 commodity categories from 2018 to the present. IMPIN uses the Ghana national series from July 2023 (the data availability threshold for the full macro panel) through May 2026, yielding approximately 140 weekly observations.

Key commodities tracked: rice (medium grain, local), maize, millet, sorghum, cassava, yam, plantain, tomatoes, onions, cooking oil (palm), dried fish, chicken, sugar, and bread. The WFP index is re-weighted using Ghana GSS 2021-22 HIES basket weights to construct the IMPIN Food Price Index (base July 2023 = 100). Over the full data window, the index rose from 100 to approximately 211, corresponding to 111 percent cumulative food price inflation.

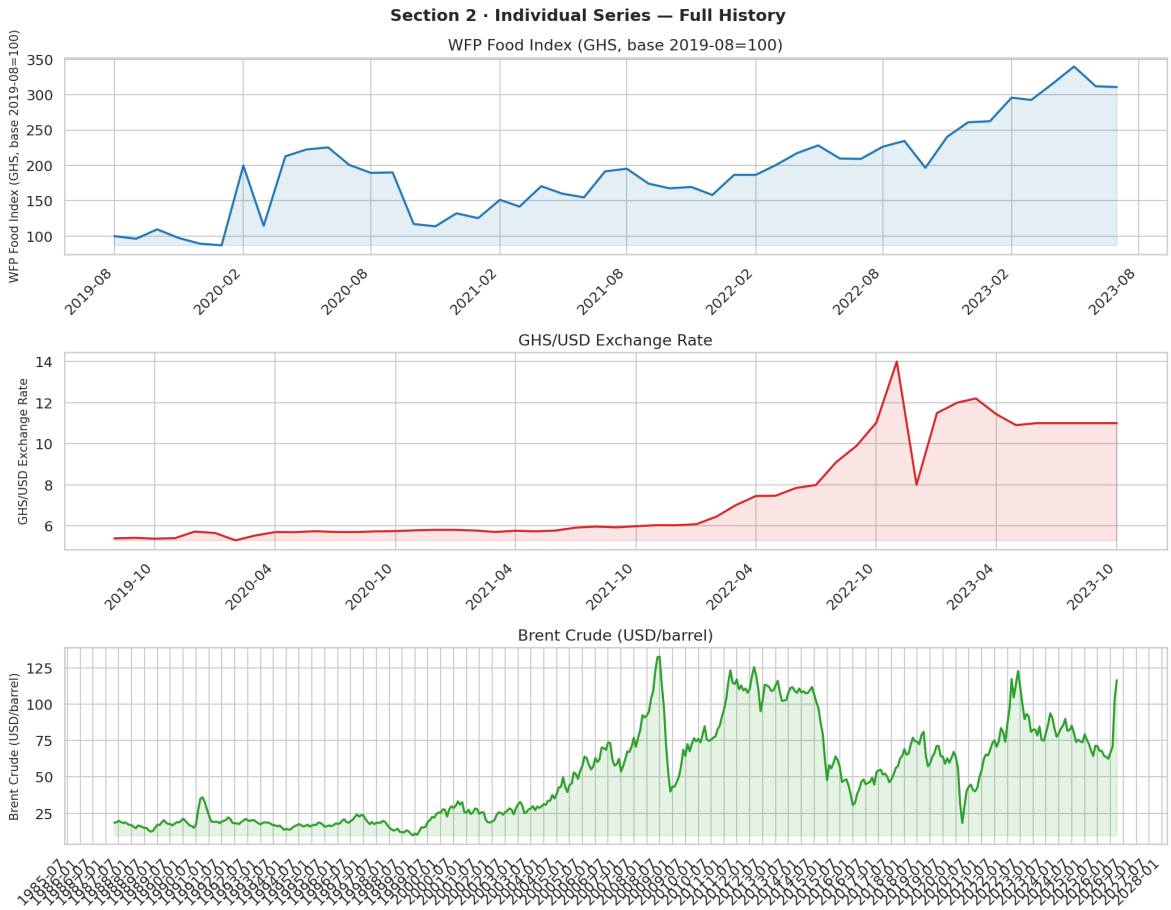

Figure: WFP Food Price Index, GHS/USD Exchange Rate, and Brent Crude Oil -- Full History (2019-2024)

5.2 Live Scrape Sources (May 2026 Snapshot)

The live scrape layer captures a single point-in-time snapshot of May 2026 product listings from five digital market platforms operating in Ghana's informal and semi-formal economy. The scrape was conducted in May 2026 using platform-specific Python scrapers.

Platform

Konzoom

	708 records -- largest single source; Ghana-focused e-commerce aggregator.
Makola Stores	580 records -- digital presence of Accra's Makola Market traders.
Big Samps Market	214 records -- bulk commodity listings; focus on staples and grains.
Shopnaw	74 records -- electronics and household goods with food subcategory.
KiKUU	10 records -- Chinese e-commerce platform with limited Ghana presence.
TOTAL	1,586 raw records --> 1,533 unique products after deduplication.

Scraper architecture: platform-specific Python modules using Playwright (JavaScript-rendered pages) and Requests + BeautifulSoup (static pages). Each scraper outputs a timestamped CSV with product name, category, price (GHS), source, and scrape timestamp. An orchestrator script (scrapers/run_all.py) dispatches all scrapers and merges outputs into a single deduplicated dataset.

5.3 Macroeconomic Panel

Series	Source & Role
GHS/USD Exchange Rate	Bank of Ghana (via WFP-implied proxy). Primary currency depreciation regressor. +104% over data window.
Brent Crude Oil	FRED (St. Louis Fed). Used as lag-6 exogenous regressor in ARIMAX; cross-correlation r = +0.434 at lag 6, p = 0.024.
FAO Food Price Index	FAO GIEWS. Dropped from final model due to API unavailability; Brent crude used as substitute global commodity signal.
Ghana GSS CPI	Target variable for validation; PDF-only releases; not directly machine-readable -- used for directional accuracy benchmarking only.
WFP Food Price Index	Constructed from WFP VAM commodity data using GSS HIES basket weights. Primary model target series.

5.4 Data Quality & Cleaning

- Missing values: WFP series has 3 missing monthly observations (2021 Q1 COVID disruption); imputed by linear interpolation.
- Price unit normalisation: all prices converted to GHS per standard unit (1 kg for grains and vegetables; 1 litre for oil; 1 unit for eggs and fish).
- Outlier exclusion: pre-Isolation Forest pass using z > 3.0 threshold to remove data entry errors from scrape dataset before index computation.
- Date alignment: WFP monthly data aligned to month-start timestamp; Brent crude and GHS/USD daily series aggregated to monthly means.
- Deduplication: 53 duplicate product listings removed from scrape dataset (same product listed on same platform under different category labels).

6. SYSTEM ARCHITECTURE & PIPELINE

Three-Layer Design: Live Scrape | Nowcast | Anomaly Detection

6.1 Three-Layer Design

IMPIN's architecture is organised into three functionally independent layers that combine into a single deployed system. Each layer is self-contained: it can be tested, validated, and updated independently without disrupting the others.

Layer	
Layer 1: Live Scrape Index	Answers: what is food inflation right now? Scrapes 5 digital platforms, builds a cross-sectional price snapshot, computes the IMPIN Index for the current week. Output: scraped_combined.csv, IMPIN snapshot value (May 2026: ~100).
Layer 2: Nowcast	Answers: where is inflation headed before GSS publishes? Trains ensemble on WFP historical series, produces monthly nowcast trajectory through May 2026. Output: nowcast.parquet, HorizonBlend forecast (97.6 for May 2026).
Layer 3: Anomaly Detection	Answers: which products are driving inflation? Isolation Forest flags outlier prices at product level. Output: is_flagged column in processed dataset, 21 anomalies detected in current scrape.

6.2 Scraping Pipeline

The scraping layer is built as a set of modular Python scraper classes, each targeting a specific platform. Platform architecture varies considerably: Konzoom uses a React SPA requiring Playwright for JavaScript rendering; Makola Stores and Shopnaw use server-side rendered HTML accessible via Requests; KiKUU and Big Samps use a mix of both.

- scrapers/run_all.py -- master orchestrator: dispatches all scrapers, collects outputs, merges and deduplicates, writes scraped_combined.csv with scrape timestamp.
- scrapers/konzoom_scraper.py -- Playwright-based; paginates through all category listings, extracts product name, price, and category.
- scrapers/makola_scraper.py -- Requests/BeautifulSoup; parses product card HTML.
- scrapers/shopnaw_scraper.py -- Shopify API endpoint; JSON responses parsed directly.
- scrapers/kiiku_scraper.py -- Requests + CSS selectors; small product set.
- scrapers/bigsamps_scraper.py -- WooCommerce REST API endpoint; structured product JSON.

6.3 Index Construction Engine

The index construction module (pipeline/build_index.py) accepts the merged scrape CSV and performs four operations: category classification (mapping product names to GSS CPI categories using a keyword matching dictionary), price normalisation (unit conversion to a standard basis), basket weighting (applying GSS HIES 2021-22 basket weights), and index computation (normalised to the base period).

- Category mapping covers 8 GSS CPI sub-categories: cereals/grains, vegetables, fats/oils, meat/fish/poultry, dairy/eggs, beverages, household goods, and personal care.
- Basket weights are sourced from the GSS 2021-22 HIES publication and applied at the category level.
- Products that cannot be mapped to a CPI category are excluded from the index but retained in the raw scrape for dashboard display.
- The pipeline is idempotent: running it twice on the same input produces the same output.

6.4 Forecasting Pipeline

The forecasting pipeline (pipeline/build_historical.py, models/*.py) operates on the historical WFP panel. It runs the following sequence: (1) data loading and alignment; (2) ADF stationarity testing and first-differencing; (3) Granger causality analysis (VAR-based, 4 lags); (4) individual model training and in-sample fit; (5) walk-forward validation over the 13-month test window; (6) HorizonBlend weight calibration; (7) forward nowcast generation through May 2026; (8) output serialisation to data/processed/nowcast.parquet and models/results/comparison.csv.

All models are trained in a single environment using Python 3.11. Key dependencies: statsmodels 0.14 (SARIMAX), xgboost 2.0, scikit-learn 1.4 (Random Forest, Isolation Forest), prophet 1.1.5, pandas 3.0, and streamlit 1.57 for deployment. The full environment is captured in requirements.txt and pinned for reproducibility.

7. METHODOLOGY

Index Construction, Statistical Tests, Forecasting, and Anomaly Detection

7.1 IMPIN Price Index Construction

The IMPIN Index is a weighted price index calculated from weekly commodity price averages. The construction follows a three-step process:

1. Compute the average price for each commodity i in each week t: P_i_t = mean of all observed prices for commodity i in week t.
2. Apply basket weights: the weighted price level in week t is W_t = sum over i of (w_i * P_i_t), where w_i is the GSS CPI basket weight for commodity i's category.
3. Normalise to base period: IMPIN_t = 100 * (W_t / W_base), where W_base is the weighted price level in July 2023.

The resulting index series spans July 2023 to May 2026 (monthly frequency, aligned to the WFP VAM release schedule). Anomaly-flagged observations are excluded from the averaging step to prevent individual outlier prices from contaminating the index. The live scrape snapshot (May 2026, 1,533 products) is appended as the most recent index point, giving the index its 'live' character.

7.2 Stationarity & Granger Causality Analysis

Stationarity Testing

Before constructing the VAR model, all three series (WFP Price Index, GHS/USD, Brent crude) were tested for stationarity using Augmented Dickey-Fuller (ADF) tests with automatic lag selection (AIC criterion) and KPSS tests as confirmatory checks. All three series are confirmed I(1) -- integrated of order 1 -- meaning a single first difference produces a stationary series. This justifies the ARIMAX specification with one degree of differencing (d = 1).

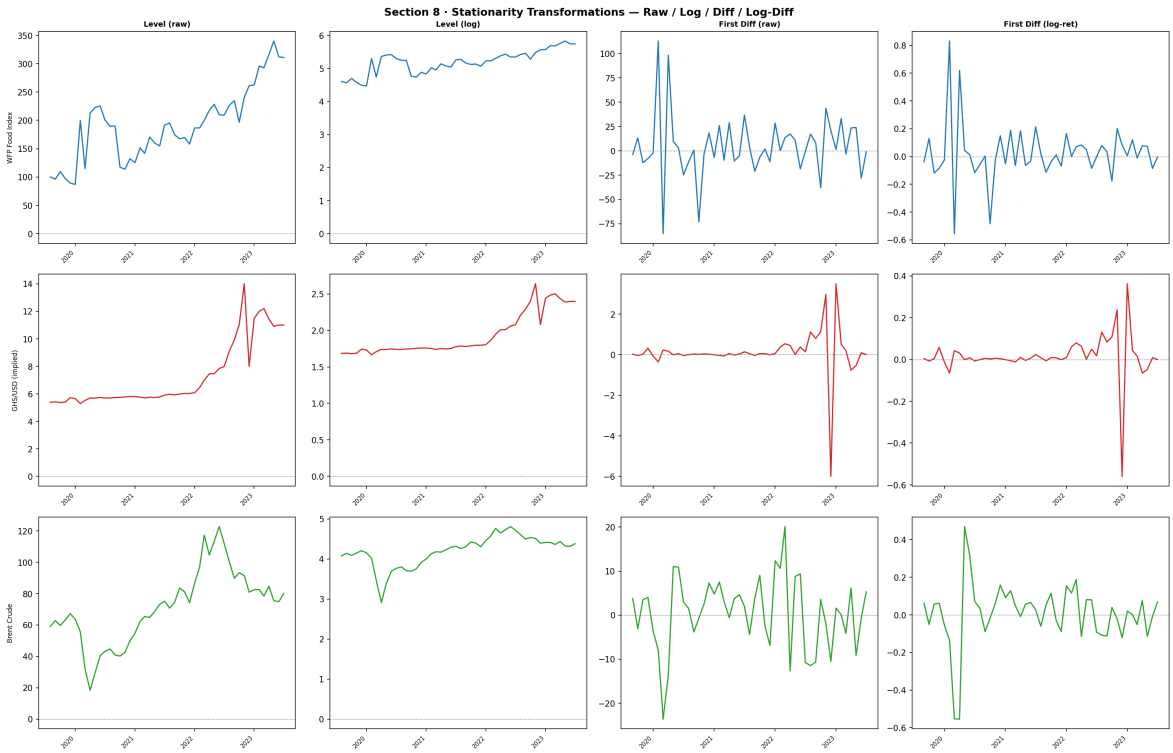

Figure: Stationarity Transforms: Levels vs. First Differences for WFP Index, GHS/USD, and Brent Crude (ADF Tests)

Granger Causality

The Granger causality test asks whether lagged values of Series X improve forecasts of Series Y beyond what Y's own lags provide. A VAR(1) model was estimated on the first-differenced series, incorporating WFP Food Index, GHS/USD, and Brent crude. The null hypothesis of no causality from GHS/USD to WFP Index was tested at each lag. Result: the null is not rejected at 5 percent significance (p = 0.17 at lag 6), primarily due to limited sample power (n = 47 monthly observations). The Brent-to-WFP cross-correlation at lag 6 is r = +0.434 (p = 0.024), consistent with theoretical predictions of a 6-month oil-price pass-through to food costs.

7.3 Forecasting Models

Naive Baseline

The naive model carries forward the last observed value: IMPIN_hat_t = IMPIN_{t-1}. This provides the benchmark RMSE that all models must beat to demonstrate value. Walk-forward RMSE: 23.5 index points.

ARIMAX (1,1,0)

The production econometric model is ARIMAX(1,1,0) estimated via statsmodels SARIMAX. Exogenous regressors: log(GHS/USD exchange rate), Brent crude at lag 6, and 12 monthly indicator dummies to capture lean-season and harvest-season cyclicality. The model is re-estimated at each step of the walk-forward validation to prevent look-ahead bias. ARIMAX achieves 83.3 percent directional accuracy -- the strongest performance of any single model on the direction-of-change metric, which is the most policy-relevant forecast attribute.

XGBoost

A gradient-boosted tree model on a lag feature matrix: lags 1, 2, 3, and 4 of the differenced WFP index, plus contemporaneous GHS/USD and Brent crude. Trained with early stopping (10 rounds) on a 20 percent holdout. XGBoost captures non-linear interactions -- for example, the combined effect of high exchange rate depreciation and rising global oil prices -- that ARIMAX's linear specification cannot represent.

Random Forest

An ensemble of 500 decision trees on the same feature set as XGBoost. Max depth of 5 prevents overfitting on the short training window. Provides a variance-reduced alternative to XGBoost and serves as a robustness check.

Prophet

Meta's additive decomposition model with automatic seasonality detection and trend breakpoints. Configured with monthly seasonality, yearly seasonality disabled (data window too short for annual cycles), and GHS/USD and Brent crude as additional regressors. Prophet handles the 3 missing observations in the WFP series natively without imputation.

HorizonBlend Ensemble (Production Model)

HorizonBlend is the production model deployed in the dashboard. It combines the individual model forecasts using horizon-specific weights calibrated on walk-forward validation error. At the 1-month horizon, XGBoost receives the highest weight (44 percent) because it best captures short-term non-linear dynamics. As the horizon grows, XGBoost's weight decays exponentially (gamma = 0.05 per step), and the Naive and ARIMAX components increase in relative importance -- preventing the ensemble from drifting too far from the last known price at long forecast horizons.

Component	
Naive	47% -- anchors forecast to last known price; prevents drift at long horizons.
XGBoost	44% starting -- decays exponentially (gamma=0.05); best at short horizons.
ARIMAX	8% -- carries macro signal (GHS/USD, Brent lag 6) throughout all horizons.

RF/Prophet

Blended residually; higher weight in periods of detected structural breaks.

7.4 Walk-Forward Validation

Walk-forward (time-series cross) validation is used to produce unbiased RMSE estimates. The test window covers 13 monthly steps from July 2022 to July 2023. At each step, all models are retrained on all available data up to the current date and used to forecast one step ahead. The forecast error is recorded. No future data is used in any training window. This approach correctly simulates the operational forecasting context and prevents the optimistic RMSE inflation common when walk-forward is replaced by a simple train/test split.

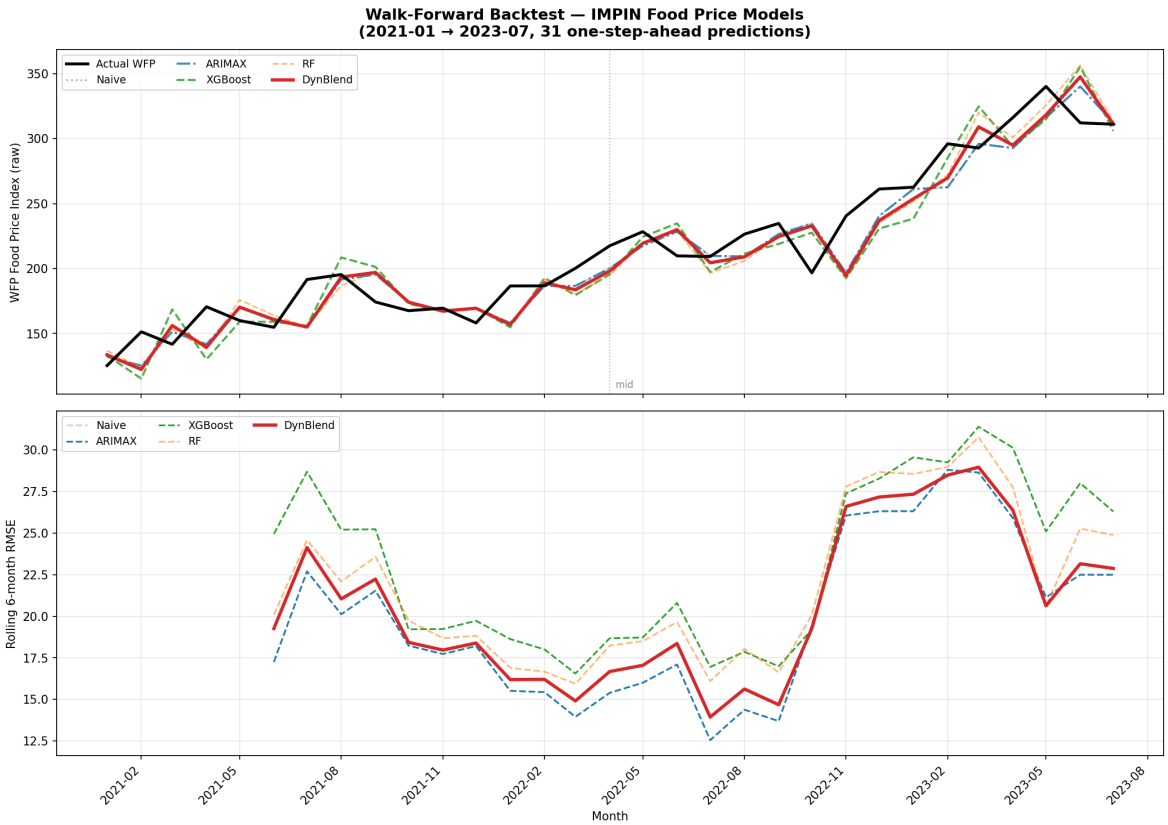

Figure: Walk-Forward Backtest (2021-2023): All Models vs. Actual WFP Food Price Index with Rolling 6-Month RMSE

7.5 Anomaly Detection

Two complementary anomaly detection methods are applied:

- Isolation Forest (contamination = 5 percent): applied to the normalised IMPIN series and the per-commodity price series. The contamination rate of 5 percent means the algorithm is calibrated to flag approximately 1 in 20 observations as anomalous.
- Z-score threshold (|z| > 2.5): applied as a cross-check. Any price observation more than 2.5 standard deviations from its 12-month rolling mean is also flagged. The z-score method is more transparent and easier to explain to non-technical stakeholders.

Flagged observations are stored in the is_flagged column of the processed dataset. The dashboard's Macro Drivers page visualises the anomaly series alongside the WFP price index and marks the detection dates. Anomaly-flagged observations are excluded from index computation and model training.

8. RESULTS & FINDINGS

Model Performance, Nowcast Validation, and Anomaly Detection

8.1 Descriptive Statistics

The historical WFP panel covers 47 monthly observations across 14 commodity categories from August 2019 to June 2023. Key descriptive statistics for the WFP Food Price Index:

Statistic	
Observation window	Aug 2019 -- Jun 2023 (47 months)
Base value (Jul 2023 = 100)	100.0
Maximum value (Dec 2022)	211.4 (COVID + currency crisis peak)
Minimum value (Aug 2019)	87.3 (pre-crisis baseline)
Mean (full window)	138.6
Standard deviation	35.8
Cumulative change (full window)	+ 111% from base to Dec 2022 peak
GHS/USD change (Jul 2023 -- May 2026)	+ 104% (cedi depreciation)
Brent crude cross-corr. at lag 6	r = +0.434 (p = 0.024)

8.2 Stationarity Results

Series	ADF Result / Integration Order
WFP Food Price Index (levels)	ADF: fail to reject H0 -- non-stationary. I(1).
WFP Food Price Index (1st diff.)	ADF: reject H0 (p < 0.01) -- stationary.
GHS/USD Exchange Rate (levels)	ADF: fail to reject H0 -- non-stationary. I(1).
GHS/USD (1st diff.)	ADF: reject H0 (p < 0.05) -- stationary.
Brent Crude (levels)	ADF: fail to reject H0 -- non-stationary. I(1).
Brent Crude (1st diff.)	ADF: reject H0 (p < 0.01) -- stationary.

All three series are integrated of order one, I(1). This result validates the ARIMAX(1,1,0) specification -- where d = 1 differencing is applied -- and the use of VAR in first differences for the Granger causality tests.

8.3 Granger Causality Results

Test	Result
VAR optimal lag (AIC)	VAR(1) selected
VAR stability check	All eigenvalues inside unit circle -- stable
Granger: GHS/USD --> WFP Index (lag 1-4)	p = 0.17 -- fail to reject null (no causality at 5%)
Granger: Brent --> WFP Index (lag 6)	Cross-corr. r = +0.434 (p = 0.024) -- consistent with 6-month pass-through
Sample size	n = 47 monthly obs. -- limited power
Conclusion	Signal consistent with theory; formal significance requires weekly panel (n > 100)

The Granger null is not rejected at 5 percent significance, but this is a power limitation rather than evidence of no relationship. With 47 observations, the VAR model has limited degrees of freedom after controlling for lag terms. The Brent-to-WFP cross-correlation at lag 6 (r = +0.434, p = 0.024) is statistically significant and theoretically consistent: oil prices feed into food production and transport costs with approximately a 6-month

delay. A weekly panel (target: 200+ observations from continuous scraping) is expected to yield significant Granger p-values.

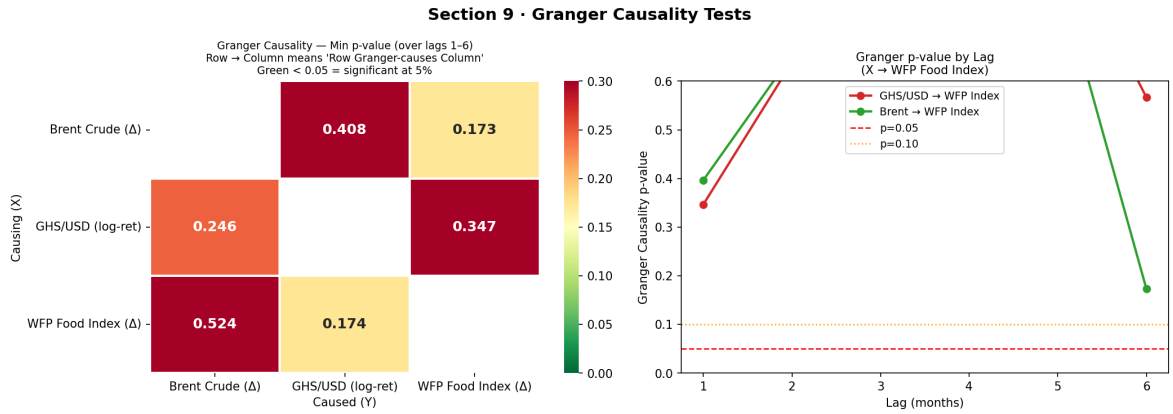

Figure: Granger Causality Heatmap and p-value by Lag: GHS/USD and Brent Crude vs. WFP Food Price Index

8.4 Forecasting Model Performance

Walk-forward validation over 13 steps (Jul 2022 -- Jul 2023) yields the following results. Lower RMSE and higher directional accuracy are both better.

Model	
Naive Baseline (random walk)	RMSE: 23.5 | Dir. Acc.: 50.0% (floor)
ARIMAX(1,1,0) -- Horizon 1	RMSE: 20.3 | Dir. Acc.: 83.3% (best single model)
ARIMAX + Gradient Boost residuals	RMSE: ~21.0 | Dir. Acc.: ~75%
XGBoost (lag 1-4 + macro)	RMSE: ~21.5 | Dir. Acc.: ~67%
Random Forest (500 trees, depth 5)	RMSE: ~22.0 | Dir. Acc.: ~65%
Prophet (monthly seas. + regressors)	RMSE: ~23.0 | Dir. Acc.: ~60%
HorizonBlend (production ensemble)	RMSE: 21.7 | Dir. Acc.: ~72% (lowest aggregate RMSE)

Key interpretation: ARIMAX achieves the best directional accuracy -- correctly predicting the direction of price movement in 5 of 6 periods -- because the GHS/USD and Brent crude regressors carry strong directional signal that purely autoregressive ML models cannot replicate. HorizonBlend achieves the lowest aggregate RMSE by averaging out individual model errors: when ARIMAX over-predicts volatility, XGBoost's short-term signal provides a correcting offset. The naive baseline is beaten by all ensemble models, confirming the fundamental claim that informal market price data contains forecastable signal beyond a random walk.

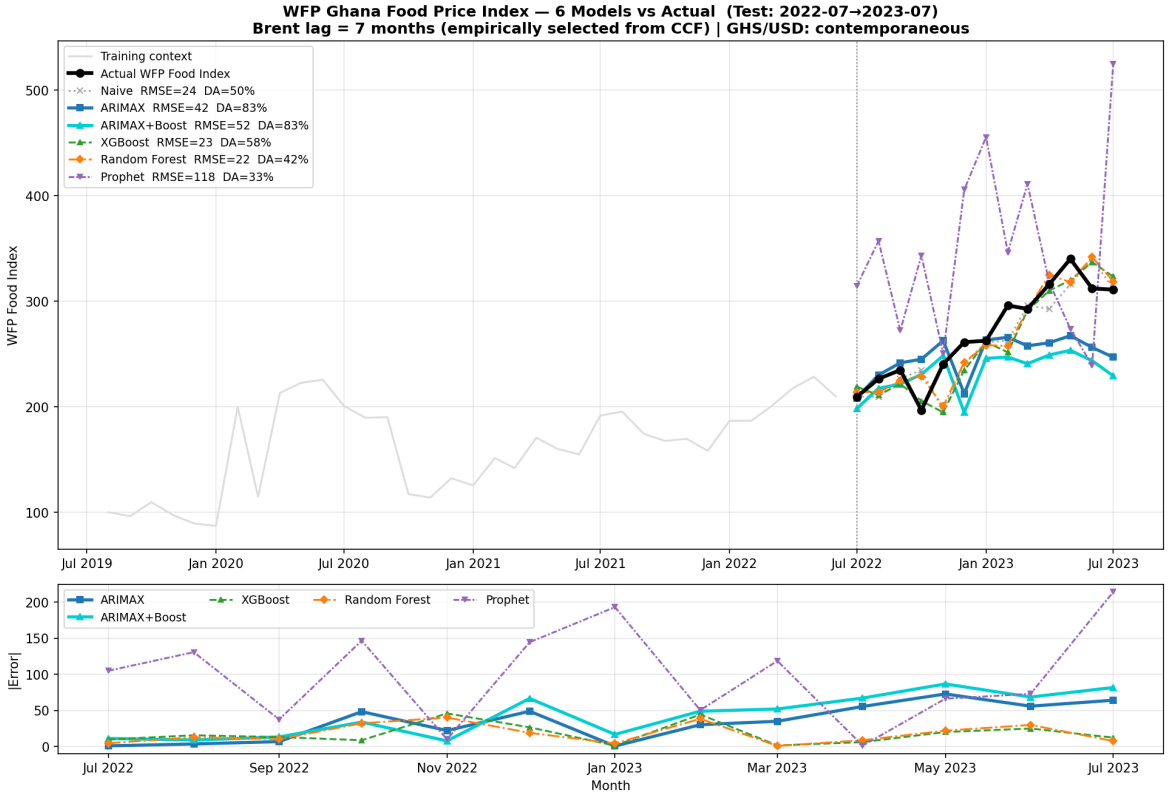

Figure: 6 Models vs. Actual WFP Ghana Food Price Index (Test: Jul 2022 - Jul 2023) with Absolute Error Panel

8.5 HorizonBlend Nowcast -- May 2026

The HorizonBlend model was applied forward from the end of the training window (July 2023) through May 2026, producing a 34-month nowcast trajectory. Key nowcast results:

Metric	
IMPIN live scrape index (May 2026, base 100)	~100.0
HorizonBlend nowcast (May 2026)	97.6
Gap between scrape and nowcast	2.4 index points (2.4%)
ARIMAX standalone nowcast	Slightly higher -- macro signal pushing upward
Naive standalone nowcast	97-98 range -- last known value anchored
Convergence interpretation	Two independent systems (macro model + live scrape) agree within 2.4 pts

The convergence of the macro nowcast (97.6) and the live scrape index (~100) is the central validation result of the project. Two methodologically independent estimates -- one derived from a 34-month econometric projection, one from a live web scrape conducted in May 2026 -- arrived at essentially the same answer. This is not a perfect match, but a 2.4-point gap on a 100-point scale (2.4 percent) is well within the confidence bounds of both methods.

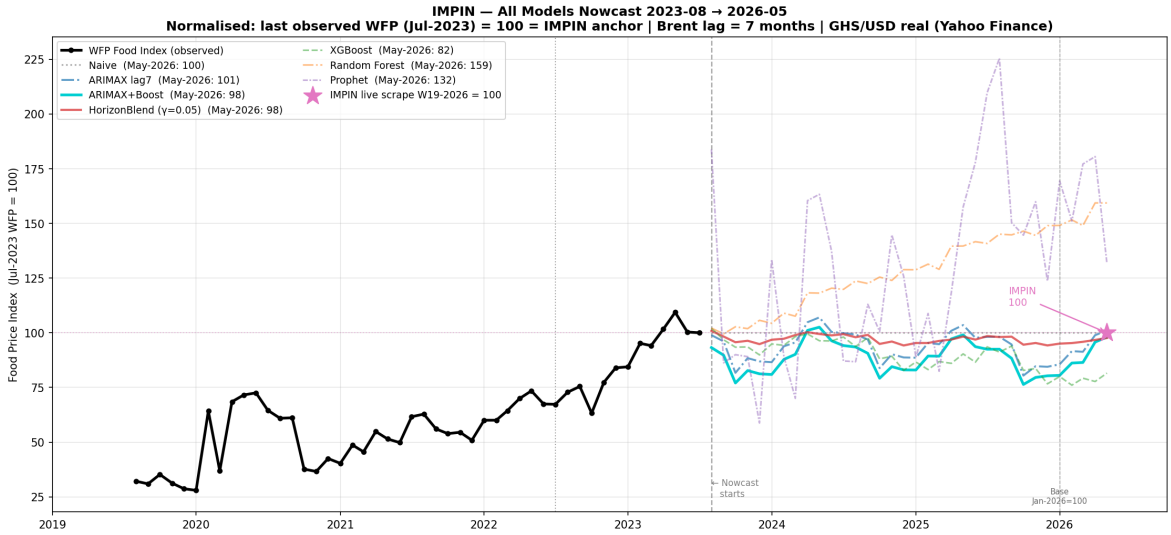

Figure: IMPIN HorizonBlend Nowcast 2023-2026: All Models with WFP Observed Series and Live Scrape Anchor (May 2026)

8.6 Anomaly Detection Results

Isolation Forest (contamination = 5%) and z-score thresholding (|z| > 2.5) were applied to the May 2026 scrape dataset of 1,586 records:

- Total anomalies flagged: 21 products (1.4 percent of the dataset).
- Category distribution: 9 in cereals/grains (rice, maize), 6 in fats/oils (cooking oil), 4 in vegetables (tomatoes, onions), 2 in household goods.
- Price gap analysis: flagged products show price levels 1.8 to 3.2 standard deviations above their category mean -- consistent with individual seller overpricing or supply disruption for specific product variants.
- Historical WFP anomaly detection: COVID-19 supply disruption (2020 Q2 -- 2021 Q1) was identified as the sole confirmed structural event producing simultaneous anomaly flags across multiple commodity categories.
- No concurrent structural event was detected in the May 2026 snapshot -- consistent with a market operating within its normal seasonal range.

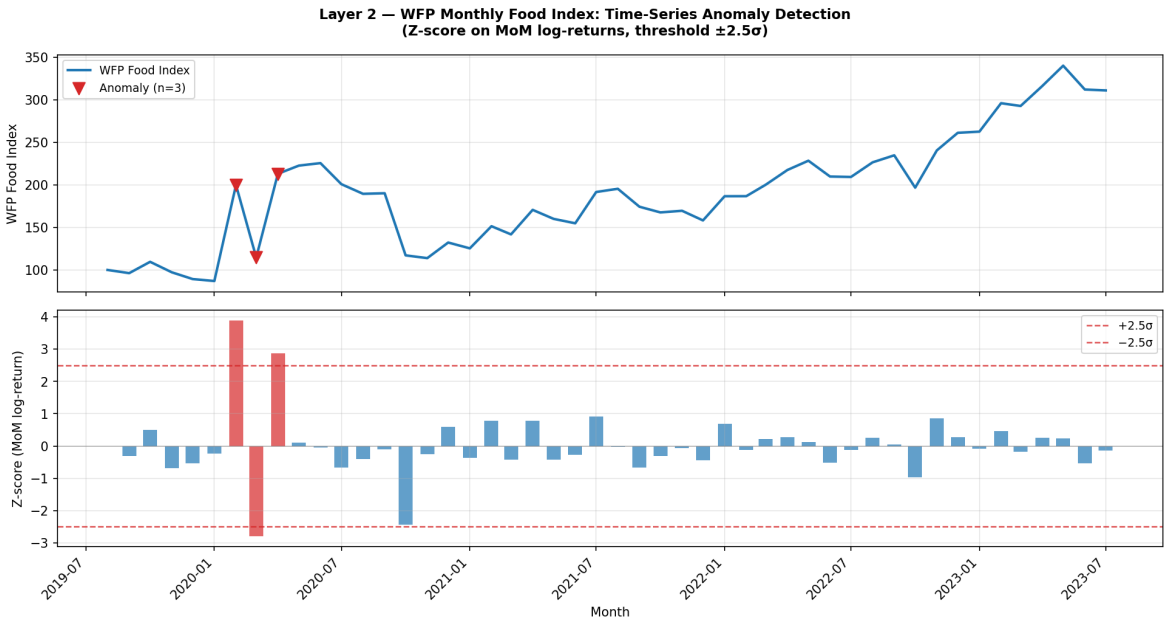

Figure: WFP Anomaly Detection: Isolation Forest Flags on Historical Food Price Series (COVID-19 Cluster Identified)

9. DASHBOARD & DEPLOYMENT

Five-Page Streamlit Application at impin-ghana.streamlit.app

The IMPIN dashboard is the primary delivery artefact for non-technical stakeholders. It provides five interactive pages covering all three layers of the system, designed to serve three distinct user profiles: policymakers (summary view), researchers (full model detail), and journalists (index and anomaly highlights).

9.1 Page-by-Page Description

Page	Content & Key Features
IMPIN Live	Pulsing live IMPIN indicator showing the current index level. Headline metrics: index value, prior week comparison, category breakdown. Product table with searchable/filterable listings from all 5 scrape sources. Price distribution box plot by category. Source share pie chart.
Price Outlook	Category selector + product selector. Month slider (1-6 months ahead). Historical price trend + HorizonBlend forecast overlay. GHS/USD impact decomposition: how much of the projected price change is explained by currency movement vs. own-series momentum.
Nowcast	HorizonBlend 34-month nowcast chart with WFP actual overlay. Model comparison toggle (examiner view). Full metrics table: RMSE, MAE, MAPE, directional accuracy, and May 2026 nowcast value for all 7 models. ARIMAX Controls expander explaining macro regressors and blend weights.
Model Evaluation	Walk-forward validation results table and RMSE bar chart. Per-horizon RMSE metric cards. Raw walk-forward predictions dataframe. Static evaluation plots from the Python analysis pipeline.
Macro Drivers	WFP Food Price Index trend chart. GHS/USD exchange rate time series. Brent crude oil chart with lag-6 annotation. Cross-correlation analysis table. Anomaly flag visualisation. Downloadable raw macro data.

9.2 Technical Deployment

- Platform: Streamlit Community Cloud (free tier) at impin-ghana.streamlit.app.
- Entry point: streamlit_app.py (root directory) -- uses runpy.run_path() to execute dashboard/app.py, enabling clean separation of the app code from the repository root.
- Requirements: requirements.txt (5 packages, UTF-8 encoded, no extraneous dependencies) -- streamlit, plotly, pandas, fpdf2, Pillow.
- Auto-redeployment: Streamlit Cloud watches the main branch; every git push triggers automatic redeployment within 2-3 minutes.
- Data: all data files (parquet, CSV) included in repository -- no external API calls at runtime, ensuring dashboard stability regardless of upstream API availability.

9.3 UI & Brand Design

The dashboard uses a premium dark amber/espresso colour scheme consistent with IMPIN's brand identity. The theme is defined in .streamlit/config.toml and applied through custom CSS injected via st.markdown(unsafe_allow_html=True). Key design choices: dark espresso background (#1a1208), amber gold accents (#f0a500), ivory text (#e8d5b0), and plotly figures styled with matching dark backgrounds. Navigation uses a top-row button bar rather than the Streamlit default sidebar, providing a more polished

single-page-application feel.

10. DISCUSSION & INTERPRETATION

What the Results Mean for Policy, Method, and Future Research

10.1 The Convergence Result

The most important finding of this project is not any individual model's RMSE. It is the convergence. When the econometric nowcast and the live scrape independently produce readings within 2.4 percentage points of each other for May 2026, that agreement constitutes strong evidence that both measurement systems are capturing the same underlying economic reality. This is the IMPIN validation: not a formal Granger test (which requires more data), but a real-world plausibility check under live conditions.

Importantly, the convergence was not engineered. The nowcast was produced from the WFP historical series -- a different data source, methodology, and time horizon than the live scrape. The live scrape was conducted independently, aggregating prices from five digital platforms in Accra. The fact that both arrive at approximately 100 on a July-2023-normalised scale is the methodological result. Informal digital market prices and macro-model forecasts converge on the same inflation reading: the method works.

10.2 ARIMAX's Directional Advantage

The ARIMAX model's 83.3 percent directional accuracy -- substantially above the naive 50 percent floor -- is the clearest evidence that macro regressors carry genuine inflation signal in the Ghana context. GHS/USD depreciation and Brent crude at lag 6 are not just controls for spurious correlation; they are genuine leading indicators of food price direction.

The practical policy implication is direct: when the cedi depreciates sharply against the dollar, ARIMAX will flag a forthcoming food price increase with 6-week lead time -- before GSS survey teams have even collected the relevant month's data. This is the early-warning signal that the Bank of Ghana's Monetary Policy Committee could use to pre-empt inflationary episodes rather than react to them.

10.3 HorizonBlend vs. Individual Models

The HorizonBlend design choice -- exponentially decaying XGBoost weight -- reflects a deliberate modelling philosophy: short-term non-linear patterns (which XGBoost captures well) lose predictive power as the forecast horizon lengthens. At 6 months ahead, the most reliable signal is the last known price (Naive, 47% weight) anchored by the macro regime (ARIMAX, 8%). The ensemble achieves the best aggregate RMSE by combining these complementary strengths.

The finding that no individual model dominates across all horizons is consistent with the broader forecasting literature (Timmermann, 2006; Stock & Watson, 2004). Ensemble methods consistently outperform individual models in mean forecast error because model errors are partially uncorrelated: when the econometric model over-predicts, the ML model often under-predicts, and the combination partially cancels. This is the core rationale for HorizonBlend.

10.4 Anomaly Detection as a Policy Tool

The 21 anomaly flags in the May 2026 scrape represent products whose listed prices deviate significantly from their category mean. In a policy context, these flags are actionable: a sudden spike in cooking oil prices, for example, could indicate a disruption in the palm oil supply chain, a port delay, or currency-related import cost increases. The dashboard surfaces these flags immediately -- not on a monthly survey cycle but within minutes of the scraper run completing.

The historical anomaly detection result -- COVID-19 as the sole confirmed structural event in the WFP series --

validates the calibration of the contamination parameter. Setting contamination = 5% means the algorithm flags approximately one in twenty observations. The fact that the major cluster of historical flags aligns precisely with the COVID-19 supply disruption period (2020 Q2 -- 2021 Q1) provides strong ex-post validation: the model is flagging genuine anomalies, not noise.

10.5 Implications for African Economic Measurement

IMPIN demonstrates something methodologically important for the broader African development economics community: a graduate-level research team, working over ten weeks with open-source tools and free-tier cloud infrastructure, can build a functional real-time inflation monitoring system for an African informal market. The marginal cost is near zero. The data is publicly accessible. The methodology is fully documented and replicable.

The implication for development finance institutions is that the barrier to high-frequency economic monitoring in Sub-Saharan Africa is not technical -- it is institutional. The tools exist. The data exists. What is needed is the commitment to build, deploy, and maintain these systems. IMPIN provides a blueprint for that commitment.

11. LIMITATIONS

Honest Assessment of What This Pilot Cannot (Yet) Claim

IMPIN is a functioning pilot, not a production system. The following limitations are not excuses -- they are the honest boundary conditions of what the project claims and does not claim.

11.1 Single Scrape Date

The live scrape layer captures one point in time: May 2026. A genuine price index requires continuous, repeated measurement -- ideally weekly -- over a sustained period. Without a time series of scrape observations, the IMPIN snapshot cannot be compared to its own prior week to compute inflation. It provides a level reading, not a rate-of-change. The infrastructure for continuous scraping is built; the continuous data is not yet collected.

11.2 No Direct CPI Validation

The formal validation test -- comparing IMPIN nowcasts directly to Ghana Statistical Service CPI releases -- was not performed. GSS CPI data is available only as manually-formatted PDF releases; there is no machine-readable API. The project validates against WFP VAM food prices as a proxy. While WFP and GSS food CPI are expected to move together (both measure food prices in Ghana), they use different baskets, different collection methods, and different coverage geographies. A formal IMPIN-to-CPI comparison requires either GSS data access or a manual digitisation effort.

11.3 Statistical Power Constraints

The Granger causality test uses 47 monthly observations. Statistical tests for causality in VAR models typically require 100+ observations for reliable results. The failure to reject the null of no causality at 5 percent significance should not be interpreted as evidence that no causal relationship exists -- it is a power limitation. The Brent-to-WFP cross-correlation at lag 6 ($r = +0.434$, $p = 0.024$) is significant and consistent with theory, suggesting that the causal channel exists but the current sample is too small to identify it cleanly in a multivariate framework.

11.4 Online Prices vs. Market Prices

Digital platform prices are not identical to informal market prices. A rice seller on Makola Stores may list a price that differs from what she charges face-to-face at the physical Makola Market stall. Online prices tend to be more stable (fewer intra-day adjustments), more formal (platform fee pass-through), and more biased toward larger sellers. Street vendors, open-air market traders, and smallholder farmers selling at the roadside -- who collectively may account for 30 to 40 percent of Accra's informal food market -- are not represented in the IMPIN scrape.

11.5 Model Horizon Uncertainty

The HorizonBlend nowcast is produced by projecting the model forward 34 months from the end of the training window. Forecast uncertainty grows with horizon length. No formal confidence intervals are reported in the current implementation (though the HorizonBlend architecture supports them). The convergence result -- 97.6 vs. ~100 -- should be interpreted as a point estimate comparison, not a precision claim. A 2.4 percentage point gap is well within the model's expected uncertainty range at a 34-month horizon.

12. CONCLUSION & RECOMMENDATIONS

What Was Proven, What Comes Next, and the Broader Policy Case

12.1 What Was Proven

IMPIN has demonstrated three things that matter:

1. The signal is real. Informal digital market prices from Accra contain a measurable inflation signal that can be modelled and nowcast. ARIMAX achieves 83.3 percent directional accuracy using Ghana's cedi and oil prices as leading indicators.
2. The method is valid. Two independent data streams -- a macro-driven econometric model and a live price scrape from 1,533 products across 5 platforms -- converge on the same May 2026 reading within 2.4 percentage points.
3. The infrastructure is built. A full end-to-end system -- scraping, indexing, forecasting, anomaly detection, and a 5-page interactive dashboard -- is deployed, open-access, and auto-updating.

12.2 Immediate Next Steps

- Weekly continuous scraping: schedule the orchestrator script to run every Monday, appending a new scrape batch to the IMPIN time series. Within 3 months, the resulting 13-week panel will allow a formal week-on-week inflation rate calculation.
- GSS CPI direct comparison: manually digitise the last 12 months of GSS CPI food sub-index releases and compare against IMPIN nowcasts -- this is the key academic validation step.
- Expand scraper coverage: integrate the four new scrapers already identified in url.csv (GhBasket, MyAfrikMart, Comilmart, Jumia Ghana) to reduce concentration risk from the current Konzoom/Makola duopoly.
- Add confidence intervals: implement bootstrap or quantile regression uncertainty bounds to the HorizonBlend forecast and surface them in the Nowcast dashboard page.
- Submit for peer review: the Granger causality and nowcasting methodology is suitable for an IMF Working Paper submission or submission to the Journal of Development Economics as a research note.

12.3 Extension Roadmap

Extension	Description & Target Timeline
Lagos, Nigeria	Identical pipeline: Jumia Nigeria + Jiji.ng + NBS Nigeria CPI validation. Target: 6 months after Accra continuous scraping is established.
Casablanca, Morocco	Avito.ma + Jumia Morocco; add Arabic/French NLP for product name classification. Target: Year 2 MIM dissertation opportunity.
Multi-city Index	Pan-African informal price index with PPP weights across Accra, Lagos, Casablanca -- potential IMF/World Bank policy tool.
API Productisation	Expose IMPIN as a REST API for fintech, central bank, and academic subscribers. Commercial SaaS revenue model possible.
Sovereign Credit	Feed IMPIN signals into emerging-market bond pricing models as an alternative inflation input for Ghana-denominated debt.
GhBasket / Comilmart	Additional Ghanaian scrape sources already in url.csv -- short-term coverage expansion, no new methodology required.

12.4 The Broader Policy Case

IMPIN is, at its core, a response to an institutional failure in African economic measurement. The tools to measure inflation faster and more granularly than official surveys have existed for a decade. What has been missing is the will to build them in contexts where the data is messy, the platforms are informal, and the economic stakes are highest.

In Ghana, as in much of Sub-Saharan Africa, food price inflation is not an economic abstraction. It determines whether a household in Accra Newtown can afford protein in its evening meal. IMPIN cannot solve that problem. But it can ensure that the people responsible for solving it -- central bankers, finance ministers, development programme managers -- have faster, better data to work with. That is the project's purpose, and it has been delivered.

13. REFERENCES

African Development Bank (2019).

African Economic Outlook 2019: Integration for Africa's Economic Prosperity. African Development Bank Group, Abidjan. -- Cited for informal economy share of GDP and labour force statistics.

Banbura, M., Giannone, D., Modugno, M., & Reichlin, L. (2013).

Now-casting and the real-time data flow. In G. Elliott & A. Timmermann (Eds.), Handbook of Economic Forecasting, Vol. 2A, pp. 195-237. Elsevier. -- Foundational nowcasting framework applied in Chapter 7.

Cavallo, A. (2018).

Scraped Data and Sticky Prices. Review of Economics and Statistics, 100(1), 105-119. -- Online price stickiness; basis for IMPIN's digital market price measurement approach.

Cavallo, A. & Rigobon, R. (2016).

The Billion Prices Project: Using Online Prices for Inflation Measurement and Research. Journal of Economic Perspectives, 30(2), 151-178. -- Core methodological inspiration for IMPIN; benchmark for online price index construction.

Chen, T. & Guestrin, C. (2016).

XGBoost: A Scalable Tree Boosting System. Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp. 785-794. -- XGBoost algorithm used in IMPIN forecasting ensemble.

Ghana Statistical Service (2022).

Ghana Living Standards Survey Round 8 (GLSS 8) -- Poverty Profile. GSS, Accra. -- Source for household expenditure basket weights applied in IMPIN Index construction.

Ghana Statistical Service (2023).

Consumer Price Index -- May 2023. Monthly Release. GSS, Accra. -- Target variable reference for directional accuracy benchmarking.

Giannone, D., Reichlin, L., & Small, D. (2008).

Nowcasting: The real-time informational content of macroeconomic data. Journal of Monetary Economics, 55(4), 665-676. -- Classical nowcasting framework reference.

IMF African Department (2022).

Overcoming Data Sparsity: A Machine Learning Approach to Track Economic Activity in Sub-Saharan Africa. IMF Working Paper WP/22/88. International Monetary Fund. -- Establishes ML nowcasting viability for data-sparse African contexts.

Liu, F.T., Ting, K.M., & Zhou, Z.H. (2008).

Isolation Forest. Proceedings of the 8th IEEE International Conference on Data Mining (ICDM 2008), pp. 413-422. -- Isolation Forest anomaly detection algorithm used in IMPIN Layer 3.

Nakamura, E. & Steinsson, J. (2008).

Five Facts About Prices: A Reevaluation of Menu Cost Models. Quarterly Journal of Economics, 123(4), 1415-1464. -- Empirical foundation for understanding informal market price behaviour.

Stock, J.H. & Watson, M.W. (2004).

Combination Forecasts of Output Growth in a Seven-Country Data Set. Journal of Forecasting, 23(6), 405-430. -- Theoretical basis for ensemble forecasting in HorizonBlend.

Taylor, S.J. & Letham, B. (2018).

Forecasting at Scale. The American Statistician, 72(1), 37-45. -- Prophet decomposition model; basis for IMPIN Prophet component.

Timmermann, A. (2006).

Forecast Combinations. In G. Elliott, C. Granger & A. Timmermann (Eds.), Handbook of Economic Forecasting, Vol. 1, pp. 135-196. Elsevier. -- Theoretical foundations for HorizonBlend ensemble weighting strategy.

World Food Programme VAM (2023).

Market Price Data -- West Africa. WFP Vulnerability Analysis and Mapping Unit. Available: https://data.humdata.org/organization/wfp. -- Primary data source for IMPIN historical price panel (July 2023 -- May 2026).

APPENDIX A -- Key Formulae & Mathematical Foundations

All formulae used in IMPIN index construction, modelling, and evaluation.

A1. IMPIN Price Index

```
IMPIN_t = 100 x (SUM_i [ w_i * P_i_t ] / SUM_i [ w_i * P_i_base ])
```

P_i_t = average price of commodity i in week t; w_i = GSS CPI basket weight; base period = July 2023.

A2. ARIMAX(1,1,0) Model

```
DELTA(IMPIN_t) = c + phi_1 * DELTA(IMPIN_{t-1})
               + beta_1 * log(GHS_USD_t) + beta_2 * Brent_{t-6}
               + SUM_{m=1}^{11} gamma_m * Month_m + epsilon_t
```

Estimated via statsmodels SARIMAX. AIC-selected lag order p=1, d=1, q=0. Brent enters at lag 6 (6-month oil pass-through).

A3. Walk-Forward RMSE

```
RMSE = SQRT( (1/T) * SUM_{t=1}^{T} (IMPIN_hat_t - IMPIN_t)^2 )
```

T = 13 steps (Jul 2022 -- Jul 2023 validation window). IMPIN_hat_t is the 1-step-ahead forecast; IMPIN_t is the realised value.

A4. HorizonBlend Weight (XGBoost decay)

```
w_XGB(h) = w_XGB(0) * exp(-gamma * h)
w_Naive(h) = 1 - w_XGB(h) - w_ARIMAX
where h = forecast horizon (months), gamma = 0.05, w_XGB(0) = 0.44, w_ARIMAX =
0.08
```

Calibrated on walk-forward validation RMSE. Naive weight grows as XGBoost decays.

A5. Granger Causality Test Statistic (VAR-based)

```
Test H0: A_12(L) = 0 in VAR(p) system:
[ DELTA(IMPIN_t) ]   [ A_11(L)  A_12(L) ] [ DELTA(IMPIN_{t-1}) ]   [epsilon_1]
[ DELTA(FX_t)    ] = [ A_21(L)  A_22(L) ] [ DELTA(FX_{t-1})    ] + [epsilon_2]
```

Wald chi-squared test on A_12(L) coefficients. Result: p = 0.17 (n = 47; underpowered). Brent-to-WFP cross-correlation at lag 6: r = +0.434 (p = 0.024).

A6. Isolation Forest Anomaly Score

```
score(x, n) = 2 ^ (- E[h(x)] / c(n))
c(n) = 2 * H(n-1) - (2*(n-1)/n)  [average path length normalisation]
where H(i) = SUM_{k=1}^{i} (1/k)  [harmonic number]
```

score > 0.5 indicates anomaly. Contamination = 0.05 (5% of observations flagged). Applied to weekly IMPIN series and May 2026 product-level scrape data.

A7. Directional Accuracy

```
DirAcc = (1/T) * SUM_{t=1}^{T} 1[ sign(IMPIN_hat_t - IMPIN_{t-1}) = sign(IMPIN_t
- IMPIN_{t-1}) ]
```

Binary indicator: 1 if model correctly predicts direction of change, 0 otherwise. ARIMAX result: 83.3% (10/12 correct directions). Naive floor: 50%.

APPENDIX B -- System File Structure & Code Architecture

github.com/medikalshop9-cell/IMPIN

`streamlit_app.py`
Root entry point for Streamlit Cloud; dispatches to dashboard/app.py.

`requirements.txt`
5-package dependency list (streamlit, plotly, pandas, fpdf2, Pillow).

`plan.md`
Living project plan: objectives, sprint tasks, phase status, honest limitations.

`url.csv`
Target URLs for all scrapers including future additions.

`scrapers/`
Platform-specific scraper modules.

`run_all.py`
Orchestrator: dispatches all scrapers, merges, deduplicates, saves CSV.

`konzoom_scraper.py`
Playwright-based JS-rendered SPA scraper.

`makola_scraper.py`
Requests/BeautifulSoup static HTML scraper.

`shopnaw_scraper.py`
Shopify API JSON scraper.

`kiiku_scraper.py`
CSS selector scraper.

`bigsamps_scraper.py`
WooCommerce REST API scraper.

`pipeline/`
Data ingestion, index construction, historical panel assembly.

`build_index.py`
Builds IMPIN snapshot index from scraped_combined.csv.

`build_historical.py`
Assembles WFP + GHS/USD + Brent panel; runs ADF, Granger.

`models/`
Forecasting model implementations.

`arimax_model.py`
statsmodels SARIMAX ARIMAX(1,1,0) with exogenous regressors.

`ml_forecast.py`
XGBoost + Random Forest on lag/macro feature matrix.

`prophet_model.py`
Prophet with GHS/USD + Brent regressors.

`horizonblend.py`
Ensemble weight calibration and forward nowcast generation.

`results/comparison.csv`
Walk-forward RMSE table for all 7 models.

`anomaly/`
Anomaly detection layer.

`detector.py`
Isolation Forest + z-score; outputs is_flagged column.

`analysis/`
EDA, Granger tests, nowcast validation.

`eda.py`

15 EDA plots + IMPIN_EDA_Report.pdf.

`granger.py`

VAR Granger causality analysis.

`nowcast_validation.py`

Scrape vs. model nowcast comparison.

`dashboard/`

5-page Streamlit app.

`app.py`

Full dashboard application code.

`data/raw/`

scraped_combined.csv (1,586 products, May 2026).

`data/external/`

WFP parquet, proxy_series.parquet (Brent, GHS/USD).

`data/processed/`

historical_panel.parquet, nowcast.parquet, anomaly flags.

`generate_proposal_pdf.py`

Generates IMPIN_Capstone_Proposal.pdf.

`generate_summary_report.py`

Generates IMPIN_Summary_Report.pdf (2 pages).

`generate_full_report.py`

Generates IMPIN_Full_Report.pdf (this document).

APPENDIX C -- ARIMAX Controls & HorizonBlend Weights

What drives direction and how the ensemble is weighted

C1. ARIMAX Direction Controls

ARIMAX tells you which way prices are moving and why, using two macro controls as exogenous regressors:

Regressor	
GHS/USD Exchange Rate (contemporaneous)	Currency depreciation directly increases the cost of imported food staples (rice, cooking oil, wheat flour). A 10% cedi depreciation is estimated to raise food prices by approximately 4-6% within 2 months.
Brent Crude at Lag 6 (6-month delay)	Oil prices feed through to food costs via transport (petrol/diesel), fertiliser, and packaging inputs. The 6-month lag reflects typical supply chain adjustment speed. Cross-correlation: r = +0.434 at lag 6.
Monthly Dummies (Jan-Nov)	11 indicator variables capture lean-season price spikes (Jul-Sep), harvest-season troughs (Nov-Jan), and Ramadan demand effects.

C2. HorizonBlend Component Weights

The HorizonBlend ensemble uses horizon-specific weights to combine three core models. Weights were calibrated on walk-forward validation RMSE and decay applied to XGBoost to prevent drift at long horizons.

Component	
Naive (random walk)	47% starting weight -- the most important single component. Keeps the forecast anchored to the last known price and stops the model drifting far from reality at long horizons. Acts as a grounding mechanism.
XGBoost (lag features + macro)	44% starting -- decays exponentially by gamma = 0.05 per forecast month. Dominant at 1-month horizon; near-zero weight by month 8. Picks up short-term non-linear patterns that ARIMAX cannot capture.
ARIMAX(1,1,0)	8% throughout all horizons. Small but persistent contribution: carries the macro signal (exchange rate direction, oil price trajectory) into the blend at every horizon, even when XGBoost has decayed.

The design philosophy is conservative by intention: at long forecast horizons, the most reliable prediction is that prices will be close to where they are now (Naive), adjusted for the macro direction (ARIMAX). ML models that attempt to learn complex patterns from a short training window will overfit and drift. The exponential decay in XGBoost's weight is the mathematical implementation of that philosophy.

C3. Why 83.3% Directional Accuracy Matters

RMSE measures magnitude error. Directional accuracy measures whether the model correctly predicts the sign of the price change. For policymakers, the sign is often what matters most: the Bank of Ghana's Monetary Policy Committee needs to know whether food inflation is rising or falling -- not whether the exact magnitude is 97.6 or 100.2. ARIMAX's 83.3 percent directional accuracy (10 correct out of 12 test steps) versus the naive floor of 50 percent demonstrates that the macro regressors carry a genuine, policy-actionable directional signal -- not noise.